

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
%matplotlib inline
```

```
In [2]: df=pd.read_csv("AAPL - AAPL.csv")
```

```
In [3]: df
```

Out[3]:

	Date	Open	High	Low	Close(t)	Volume	SD20	Upper_Band	Lower_Ban
0	2005-10-17	6.66	6.69	6.50	6.60	154208600	0.169237	6.827473	6.15052
1	2005-10-18	6.57	6.66	6.44	6.45	152397000	0.168339	6.819677	6.14632
2	2005-10-19	6.43	6.78	6.32	6.78	252170800	0.180306	6.861112	6.13988
3	2005-10-20	6.72	6.97	6.71	6.93	339440500	0.202674	6.931847	6.12115
4	2005-10-21	7.02	7.03	6.83	6.87	199181500	0.216680	6.974860	6.10814
...	...	...	...	...	...	...	...	...	...
3727	2020-08-07	452.82	454.70	441.17	444.45	49453300	27.954399	455.316298	343.49870
3728	2020-08-10	450.40	455.10	440.00	450.91	53100900	29.847338	462.586675	343.19732
3729	2020-08-11	447.88	449.93	436.43	437.50	46975600	30.576290	466.543079	344.23792
3730	2020-08-12	441.99	453.10	441.19	452.04	41486200	32.050532	472.583564	344.38143
3731	2020-08-13	457.72	464.17	455.71	460.04	52520500	33.532634	479.279768	345.14923

3732 rows × 64 columns

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3732 entries, 0 to 3731
```

```
Data columns (total 64 columns):
```

#	Column	Non-Null Count	Dtype
0	Date	3732 non-null	object
1	Open	3732 non-null	float64
2	High	3732 non-null	float64
3	Low	3732 non-null	float64
4	Close(t)	3732 non-null	float64
5	Volume	3732 non-null	int64
6	SD20	3732 non-null	float64
7	Upper_Band	3732 non-null	float64
8	Lower_Band	3732 non-null	float64
9	S_Close(t-1)	3732 non-null	float64
10	S_Close(t-2)	3732 non-null	float64
11	S_Close(t-3)	3732 non-null	float64
12	S_Close(t-5)	3732 non-null	float64
13	S_Open(t-1)	3732 non-null	float64
14	MA5	3732 non-null	float64
15	MA10	3732 non-null	float64
16	MA20	3732 non-null	float64
17	MA50	3732 non-null	float64
18	MA200	3732 non-null	float64
19	EMA10	3732 non-null	float64
20	EMA20	3732 non-null	float64
21	EMA50	3732 non-null	float64
22	EMA100	3732 non-null	float64
23	EMA200	3732 non-null	float64
24	MACD	3732 non-null	float64
25	MACD_EMA	3732 non-null	float64
26	ATR	3732 non-null	float64
27	ADX	3732 non-null	float64
28	CCI	3732 non-null	float64
29	ROC	3732 non-null	float64
30	RSI	3732 non-null	float64
31	William%R	3732 non-null	float64
32	SO%K	3732 non-null	float64
33	STD5	3732 non-null	float64
34	ForceIndex1	3732 non-null	int64
35	ForceIndex20	3732 non-null	int64
36	Date_col	3732 non-null	object
37	Day	3732 non-null	int64
38	DayofWeek	3732 non-null	int64
39	DayofYear	3732 non-null	int64
40	Week	3732 non-null	int64
41	Is_month_end	3732 non-null	int64
42	Is_month_start	3732 non-null	int64
43	Is_quarter_end	3732 non-null	int64
44	Is_quarter_start	3732 non-null	int64
45	Is_year_end	3732 non-null	int64
46	Is_year_start	3732 non-null	int64
47	Is_leap_year	3732 non-null	int64
48	Year	3732 non-null	int64
49	Month	3732 non-null	int64
50	QQQ_Close	3732 non-null	float64
51	QQQ(t-1)	3732 non-null	float64
52	QQQ(t-2)	3732 non-null	float64
53	QQQ(t-5)	3732 non-null	float64
54	QQQ_MA10	3732 non-null	float64
55	QQQ_MA20	3732 non-null	float64

```
56  QQQ_MA50          3732 non-null    float64
57  SnP_Close         3732 non-null    float64
58  SnP(t-1))         3732 non-null    float64
59  SnP(t-5)          3732 non-null    float64
60  DJIA_Close        3732 non-null    float64
61  DJIA(t-1))        3732 non-null    float64
62  DJIA(t-5)         3732 non-null    float64
63  Close_forecast    3732 non-null    float64
dtypes: float64(46), int64(16), object(2)
memory usage: 1.8+ MB
```

In [5]:

df = df.rename(columns={'Close(t)': 'Close'})
df.head()

Out[5]:

	Date	Open	High	Low	Close	Volume	SD20	Upper_Band	Lower_Band	S_Close
0	2005-10-17	6.66	6.69	6.50	6.60	154208600	0.169237	6.827473	6.150527	6.60
1	2005-10-18	6.57	6.66	6.44	6.45	152397000	0.168339	6.819677	6.146323	6.45
2	2005-10-19	6.43	6.78	6.32	6.78	252170800	0.180306	6.861112	6.139888	6.78
3	2005-10-20	6.72	6.97	6.71	6.93	339440500	0.202674	6.931847	6.121153	6.93
4	2005-10-21	7.02	7.03	6.83	6.87	199181500	0.216680	6.974860	6.108140	6.87

5 rows × 64 columns

In [6]:

df.shape

Out[6]: (3732, 64)


In [7]:

df.columns

Out[7]: Index(['Date', 'Open', 'High', 'Low', 'Close', 'Volume', 'SD20', 'Upper\_Band', 'Lower\_Band', 'S\_Close(t-1)', 'S\_Close(t-2)', 'S\_Close(t-3)', 'S\_Close(t-5)', 'S\_Open(t-1)', 'MA5', 'MA10', 'MA20', 'MA50', 'MA200', 'EMA10', 'EMA20', 'EMA50', 'EMA100', 'EMA200', 'MACD', 'MACD\_EMA', 'ATR', 'ADX', 'CCI', 'ROC', 'RSI', 'William%R', 'SO%K', 'STD5', 'ForceIndex1', 'ForceIndex20', 'Date\_col', 'Day', 'DayofWeek', 'DayofYear', 'Week', 'Is\_month\_end', 'Is\_month\_start', 'Is\_quarter\_end', 'Is\_quarter\_start', 'Is\_year\_end', 'Is\_year\_start', 'Is\_leap\_year', 'Year', 'Month', 'QQQ\_Close', 'QQQ(t-1)', 'QQQ(t-2)', 'QQQ(t-5)', 'QQQ\_MA10', 'QQQ\_MA20', 'QQQ\_MA50', 'SnP\_Close', 'SnP(t-1))', 'SnP(t-5)', 'DJIA\_Close', 'DJIA(t-1))', 'DJIA(t-5)', 'Close\_forecast(t)'], dtype='object')

```
In [9]: df['Close'].plot()
```

Out[9]: <Axes: >




```
In [27]: df=df.drop(columns='Date')
```

```
In [28]: df.isna()
```

Out[28]:

	Open	High	Low	Close	Volume	SD20	Upper_Band	Lower_Band	S_Close(t-1)	S_Cl
0	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	
...	...	...	...	...	...	...	...	...	...	
3727	False	False	False	False	False	False	False	False	False	
3728	False	False	False	False	False	False	False	False	False	
3729	False	False	False	False	False	False	False	False	False	
3730	False	False	False	False	False	False	False	False	False	
3731	False	False	False	False	False	False	False	False	False	

3732 rows × 62 columns



In [29]: `df.isna().sum()`

```
Out[29]: Open          0
         High          0
         Low           0
         Close         0
         Volume        0
         ..
         SnP(t-5)      0
         DJIA_Close    0
         DJIA(t-1))    0
         DJIA(t-5)     0
         Close_forecast 0
         Length: 62, dtype: int64
```

In [30]: `x=df.drop(columns='Close_forecast')`

In [31]: `y=df['Close_forecast']`

In [32]: `xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2)`

In [33]: `lr=LinearRegression()`  
`lr.fit(xtrain,ytrain)`

```
Out[33]: ▼ LinearRegression
         LinearRegression()
```

In [34]: `lr.coef_`

```
Out[34]: array([-8.02951549e-02,  1.41958892e-01,  2.58713797e-01, -1.38107297e+00,
                1.24160350e-09, -1.27127620e-02, -1.49770889e-01, -9.88768867e-02,
               -1.40576003e+00, -7.08865561e-01, -4.47781386e-01, -4.43556261e-01,
               -3.33762401e-02, -1.15586007e+00, -1.11997874e+00, -1.24321467e-01,
                1.51235905e-01,  2.82769344e-02,  1.51549189e+00,  1.51549189e+00,
                1.51549189e+00,  1.51549189e+00,  1.51549189e+00, -4.46336162e-01,
                8.87577147e-01,  1.87147289e-01,  1.66144853e-03, -1.02756692e-12,
               -1.89413945e-02,  2.11753316e-03,  4.42909211e-03,  4.42844138e-03,
               -1.14148485e+01, -1.10941389e-09, -2.54720689e-11, -1.41276944e-02,
               -1.81023783e-02,  1.01553173e-02, -2.63428660e-03,  4.26103903e-01,
               -2.12586360e-01, -6.97573709e-01, -8.04242246e-02,  2.21361179e+00,
               -1.08135723e-13, -1.04435500e-01, -4.68719705e-02, -3.02189742e-01,
               -1.70048633e-01,  1.01977536e-01, -1.81289986e-03, -2.96940462e-02,
               -2.74764154e-02,  1.03770778e-01,  3.82482039e-02,  4.61062191e-03,
                2.54734444e-04, -2.82898298e-03, -1.07284925e-03,  6.24776181e-04,
                5.76527133e-05])
```

In [35]: `lr.intercept_`

```
Out[35]: 95.57516536347127
```

In [36]: `1 lr.score(xtrain,ytrain)`

```
Out[36]: 0.9991247011916671
```

```
In [39]: ypred=lr.predict(xtest)
```

```
In [40]: metrics.mean_absolute_error(ytest,ypred)
```

```
Out[40]: 1.2437301966724754
```

```
In [41]: metrics.r2_score(ytest,ypred)
```

```
Out[41]: 0.9989243375699739
```

```
In [42]: metrics.mean_squared_error(ytest,ypred)
```

```
Out[42]: 6.812385609083958
```

```
In [46]: df1=pd.DataFrame(ytest.values,columns=['actual'])
```

```
In [47]: df1
```

```
Out[47]:
```

	actual
0	55.55
1	155.52
2	11.23
3	11.15
4	163.13
...	...
742	43.02
743	15.30
744	17.53
745	72.48
746	48.00

747 rows × 1 columns

```
In [48]: df1['predicted']=ypred
```

```
In [49]: df1
```


Out[49]:

	actual	predicted
0	55.55	52.730956
1	155.52	155.915093
2	11.23	10.598524
3	11.15	10.976693
4	163.13	164.460614
...	...	...
742	43.02	42.928119
743	15.30	15.977212
744	17.53	17.738846
745	72.48	72.423698
746	48.00	46.709091

747 rows × 2 columns

```
In [51]: df1[['actual','predicted']].plot()
```


Out[51]: <Axes: >





```
In [53]: df1[['actual', 'predicted']].plot()
```

```
Out[53]: <Axes: >
```



```
In [ ]:
```